

LaRT v2.10 User Manual

Contents

1. Introduction	3
2. Installation and Build	3
2.1 Prerequisites	3
2.2 Compilation	4
2.3 Running	4
3. Input File Format	4
4. Grid Configuration	5
4.1 Cartesian Grid	5
4.2 Grid Geometry	5
4.3 Symmetry and Periodic Boundaries	5
4.4 Distance Units	6
4.5 Grid Type Selection	6
5. Source Configuration	6
5.1 Source Geometry	6
5.2 Spectral Type	7
5.3 Line Identification	7
5.3.1 Combined Hydrogen + Deuterium Ly α	8
5.3.2 He I 10833 Coherent Triplet Treatment	9
5.4 Frequency and Wavelength Grid	10
6. Medium Properties	11
6.1 Optical Depth Normalization	11
6.2 Temperature	11
6.3 Density	12
6.4 Velocity Fields	12
7. AMR Grid Mode	14
7.1 Enabling AMR Mode	14
7.2 Generic AMR Data Format	14
7.3 Extended AMR Physics Parameters	15
7.4 Tau Normalization Convention	15
7.5 Converting RAMSES Data	16
7.6 Converting Illustris/IllustrisTNG Data	16

7.7	Source Geometries in AMR Mode	18
8.	Clumpy Medium Mode	18
8.1	Overview	18
8.2	Enabling Clump Mode	18
8.3	Clump Parameters	19
8.4	Spatially-varying clump properties	20
8.5	Standalone clump generator (make_clumps.x)	21
8.6	Overlapping Clumps	22
8.7	Clump Velocity	23
9.	Observers and Peel-off	23
9.1	Observer Placement	24
9.2	Multiple Observers	24
9.3	Image Parameters	24
9.4	Peel-off Output Flags	24
9.5	Sight-line Optical-depth and Column-density Maps	25
9.6	Inside Observer (HEALPix All-Sky)	26
10.	Output Files	27
10.1	Primary Output Arrays	27
10.2	Angular Spectrum Jmu	28
10.3	Peel-off cube units and conversion to Jmu units	28
10.4	Output Control Parameters	30
11.	File Formats: FITS and HDF5	30
11.1	Format selection	30
11.2	HDF5 schema	31
11.3	Conversion utility	32
11.4	HDF5 library requirements	32
12.	Dust and Polarization	32
13.	Parallelism	32
14.	Complete Parameter Reference	33
15.	Python Utilities	37
15.1	Loading and analysing output: read_lart.py	37
15.2	Format-agnostic reader and FITS ↔ HDF5 converter: lart_io.py	40

16. Example Input Files	41
16.1 Uniform Sphere, $\text{Ly}\alpha$, Static	41
16.2 Sphere with Hubble Outflow and Peel-off	41
16.3 Realistic Galaxy (FITS Density Field, Stokes Polarization)	42
16.4 AMR Sphere (Generic Text Format)	42
16.5 Clumpy Sphere with Random Velocities	43
16.6 Illustris/TNG Galaxy (Diffuse Emissivity)	43

1. Introduction

LaRT (**L**ayman-**a**lpha **R**adiative **T**ransfer) is a Monte Carlo radiative transfer code for computing the propagation of resonance-line photons through astrophysical neutral-gas media. The default transition is Lyman- α ($\text{Ly}\alpha$, $\text{H I } 1216 \text{ \AA}$), but many other UV/optical/IR lines are supported (see §5.3). Dust scattering and absorption, full Stokes polarization tracking, Hubble-flow and other bulk velocity fields, and spatially varying density, temperature, and velocity from external data files are all supported.

Starting from v2.00, the Cartesian grid is supplemented by:

- **AMR mode** — octree adaptive mesh refinement, allowing spatially varying resolution (§7.).
- **Clumpy medium mode** — a population of compact, optically thick spherical clumps embedded in a transparent inter-clump medium (§8.).

This manual covers v2.00 and v2.10. The only user-visible change in v2.10 is the introduction of `par%grid_type` as a unified selector for the three grid modes; the old `par%use_amr_grid` and `par%use_clump_medium` flags remain valid for backward compatibility.

2. Installation and Build

2.1 Prerequisites

- Fortran 2003+ compiler: Intel `ifort/ifx` or GNU `gfortran` ≥ 7 .
- CFITSIO library.
- HDF5 library (Fortran bindings), version ≥ 1.14 , built with the same compiler used for LaRT. Required for the default build (`HDF5=1`); skip with `HDF5=0` for a FITS-only build.
- MPI library (MPICH or OpenMPI).

2.2 Compilation

```
cd LaRT_v2.00          # or LaRT_v2.10
make                   # produces LaRT.x (HDF5=1 is the default)
make HDF5=0            # FITS-only build (no HDF5 link dependency)
make HDF5_PREFIX=/path # override HDF5 install location
make DEBUG=1           # full debug flags + runtime checks
make F90=mpif90        # override compiler
make clean             # remove .o and .mod files
make cleanall          # also remove executables
make convert_ramses    # auxiliary RAMSES converter
```

Every make invocation recompiles all objects (no incremental builds). `par%HDF5_PREFIX` defaults to `/data/opt/hdf5_intel` on this machine; set it on the command line or as an environment variable when building elsewhere.

2.3 Running

```
mpirun -np 8 LaRT.x input.in
```

Output is written to `<base>.h5` by default (or `<base>.fits.gz` when `par%file_format = 'fits'` is set), where `<base>` is the input file base name (without `.in`) unless `par%out_file` is set explicitly. See §11. for the two supported file formats and how to choose between them.

3. Input File Format

Input files use Fortran *namelist* syntax. All parameters are prefixed with `par%` and collected in a single `¶meters` block.

```
1 &parameters
2   par%no_photons      = 1e6
3   par%temperature     = 1.0e4      ! gas temperature [K]
4   par%taumax         = 1.0e4
5   par%nx = 101, par%ny = 101, par%nz = 101
6   par%rmax           = 1.0
7   par%source_geometry = 'point'
8   par%spectral_type   = 'voigt'
9   par%nxfreq         = 121
10 /
```

- String values must be enclosed in single quotes.
- Multiple parameters may appear on one line, separated by commas.

- Lines beginning with ! are comments.
- Unspecified parameters retain compiled defaults (listed in §14.).

4. Grid Configuration

4.1 Cartesian Grid

The simulation domain is a rectangular box centered at the origin. By default $x \in [-x_{\max}, +x_{\max}]$ (and similarly for y and z). The cell size is $\Delta x = 2x_{\max}/n_x$.

Parameter	Default	Description
par%nx, par%ny, par%nz	101	Number of cells per axis
par%xmax, par%ymax, par%zmax	1.0	Half-box size per axis (code units)
par%rmax	−999	Spherical boundary: density = 0 for $r > r_{\max}$
par%rmin	−999	Inner cavity: density = 0 for $r < r_{\min}$
par%cone_opening	0.0	Biconical outflow half-opening angle [degrees] along the z -axis. When > 0 , density is 0.

If par%input_field or par%dens_file is given, (nx, ny, nz) and (xmax, ymax, zmax) are read from the FITS header automatically.

4.2 Grid Geometry

par%geometry selects the spatial symmetry used for accumulating and saving the mean-intensity array J :

Value	J output geometry
'sphere' (default)	1-D radial profile $J(r)$
'cylinder'	2-D cylindrical $J(\rho, z)$
'rectangle'	3-D Cartesian cube $J(x, y, z)$
'plane_atmosphere'	1-D slab $J(z)$
'spherical_atmosphere'	1-D spherical shell $J(r)$

4.3 Symmetry and Periodic Boundaries

Parameter	Default	Effect
par%xyz_symmetry	.false.	Simulate only the $+x, +y, +z$ octant ($8\times$ RAM saving)
par%xy_symmetry	.false.	Simulate only the $+x, +y$ quadrant ($4\times$ RAM saving)
par%xy_periodic	.false.	Periodic in x, y : photons exiting one face re-enter the opposite

Symmetry options are disabled automatically when peel-off is active.

4.4 Distance Units

Parameter	Default	Description
par%distance_unit	''	Physical unit: 'kpc', 'pc', 'au', or '' (dimensionless)
par%distance2cm	1.0	Alternatively: 1 code unit in cm

In dimensionless mode (par%distance_unit = ''), the gas density is determined entirely by par%taumax and the grid geometry; no physical unit is needed.

4.5 Grid Type Selection

Parameter	Default	Description
par%grid_type [†]	'car'	v2.10+ : 'car' (Cartesian), 'amr', or 'clump'
par%use_amr_grid	.false.	v2.00 : enable AMR (deprecated in v2.10)
par%use_clump_medium	.false.	v2.00 : enable clump mode (deprecated in v2.10)

5. Source Configuration

5.1 Source Geometry

par%source_geometry controls the spatial distribution of photon emission:

Value	Description
'point'	Point source at (par%xs_point, par%ys_point, par%zs_point); default (0,0,0)
'uniform'	Uniform volume emissivity throughout the grid (or inside par%source_rmax)
'uniform_sphere'	Uniform within a sphere of radius par%source_rmax
'uniform_xy'	Uniform in the $z = 0$ plane
'gaussian'	Gaussian: $\exp(-r^2/2\sigma_r^2)\exp(-z^2/2\sigma_z^2)$; widths par%source_rscale, par%source_zscale
'exponential'	Exponential disc: $\exp(-r/r_s)\exp(- z /z_s)$; scales par%source_rscale, par%source_zscale
'exponential_sphere'	Exponential clipped to sphere of radius par%source_rmax
'exponential_cylinder'	Exponential clipped to a cylinder
'sersic' / 'ssh'	Sérsic profile; index par%sersic_m, half-light radius par%Reff
'diffuse_emissivity'	3-D emissivity read from par%emiss_file
'star_file'	Individual point sources from text file par%star_file (columns: x, y, z, L)

Associated parameters:

Parameter	Default	Description
par%xs_point, par%ys_point, par%zs_point	0.0	Point source position (code units)
par%source_rmax	−999	Outer truncation radius for extended sources
par%source_rscale	−999	Radial scale length
par%source_zscale	0.0	Vertical scale height
par%sersic_m	4.0	Sérsic index
par%Reff	1.0	Effective (half-light) radius
par%comoving_source	.true.	Source photons co-moving with the local gas

5.2 Spectral Type

par%spectral_type sets the frequency distribution of emitted photons:

Value	Description
'monochromatic'	Delta function at line center ($x = 0$)
'voigt'	Voigt profile at the local gas temperature
'voigt0'	Voigt profile at the fixed temperature par%temperature0
'gaussian'	Gaussian with $\sigma = \text{par\%gaussian_sigma_vel}$ or FWHM=par%gaussian_FWHM_vel [km s ^{−1}]
'continuum'	Flat spectrum over [par%velocity_min, par%velocity_max]
'continuum+gaussian'	Fraction f_{line} as Gaussian (FWHM=par%gaussian_FWHM_vel), rest flat continuum; $f_{\text{line}} = W_{\lambda}^v / (W_{\lambda}^v + \Delta v)$ where $W_{\lambda}^v = \text{EW} / \lambda_0 c$ and Δv is the velocity range. Requires par%EW_line and par%gaussian_FWHM_vel.
'line_prof_file'	Arbitrary profile read from par%line_prof_file

5.3 Line Identification

The atomic transition is selected with par%line_id (default: Ly α).

par%line_id	Transition	Wavelength
(default)	H I Ly α	1215.67 Å
'ly_alpha_HD'	H I + D I Ly α	1215.67 Å (see below)
'CII_1334'	C II (+ fluorescence)	1334.53 Å
'CIV_1548'	C IV doublet	1548.19 / 1550.77 Å
'NV_1239'	N V doublet	1238.82 / 1242.80 Å
'OVI_1032'	O VI doublet	1031.91 / 1037.61 Å
'SiIV_1394'	Si IV doublet	1393.76 / 1402.77 Å
'NaI_D'	Na I D doublet	5891.58 / 5897.56 Å
'CaII_HK'	Ca II H & K doublet	3934.78 / 3969.59 Å
'MgII_2796'	Mg II doublet	2796.35 / 2803.53 Å
'AlII_1671'	Al II singlet	1670.79 Å
'SiII_1190' / 'SiII_1193'	Si II doublet (+ fluorescence)	1190.42 / 1193.29 Å
'SiII_1260'	Si II (+ fluorescence)	1260.42 Å
'SiII_1304'	Si II (+ fluorescence)	1304.37 Å
'SiII_1527'	Si II (+ fluorescence)	1526.71 Å
'FeII_2250'	Fe II (+ fluorescence)	2249.88 Å
'FeII_2261'	Fe II (+ fluorescence)	2260.78 Å
'FeII_UV3' / 'FeII_2344'	Fe II UV3 multiplet	2344.21 Å
'FeII_UV2' / 'FeII_2383'	Fe II UV2 multiplet	2382.76 Å
'FeII_UV1' / 'FeII_2600'	Fe II UV1 multiplet	2600.17 Å
'HeI_10833'	He I triplet	10832.06 Å

The numerical line%line_type that controls the dispatch of par%calc_voigt / par%do_resonance (set internally by setup_resonance_line) takes the following values: 1 = singlet (one upward transition); 2 = doublet (two upward transitions sharing a single thermal width, e.g. C IV, N V, O VI, Na I D, Ca II H&K, Mg II, Si IV); 4 = one resonance plus one or more fluorescent decays (most Si II, Fe II UV3); 5 = multiple upward + multiple downward transitions, the most general case (Si II 1190/1193, Fe II UV1, Fe II UV2); 6 = three upward transitions plus one downward transition (He I 10833); 7 = combined H + D Ly α (ly_alpha_HD, two coexisting two-level resonance scatterers, no fine structure).

5.3.1 Combined Hydrogen + Deuterium Ly α

The line id 'ly_alpha_HD' treats neutral hydrogen and deuterium Ly α as two coexisting two-level resonance scatterers in the same medium. The two transitions sit only 0.33 Å (~ 82 km/s) apart and have nearly identical atomic constants except for the atomic mass and rest wavelength. The deuterium-to-hydrogen number-abundance ratio is set by

Parameter	Default	Description
par%D_to_H_ratio	1.5×10^{-5}	n_D/n_H (cosmic primordial)
par%include_deuterium	.false.	Convenience flag

Setting `par%include_deuterium = .true.` together with `par%line_id = 'ly_alpha'` promotes the run automatically to `'ly_alpha_HD'`, so existing input files can opt in without changing the line id.

Physical assumptions in this version (v1):

- Both H and D are treated as simple two-level resonance scatterers (`line_type=7`). Fine-structure splitting is *not* included for either species; if `par%fine_structure = .true.` is set together with `'ly_alpha_HD'`, the code emits a warning and disables fine structure.
- The phase function is the simple dipole ($E_1 = 1, E_2 = 0, E_3 = 1$), identical for both species.
- Oscillator strengths and damping rates are taken to be equal for H and D ($f_D = f_H, \Gamma_D = \Gamma_H$ to better than 0.1%). Mass and rest wavelength differences are fully accounted for in the Voigt evaluation and species selection.
- Photon `xfreq` is carried in H-frame Doppler units throughout. D-frame conversions happen on-the-fly in the combined Voigt and scattering routines via constants precomputed at line setup (`line%delta_nu_HD_Hz`, `line%ratio_Dfreq_HD`, `line%ratio_voigta_HD`). No per-cell deuterium arrays are stored.
- During a deuterium scattering event, the perpendicular thermal velocity components and the recoil shift are evaluated using the H atomic constants, which produces a small ($\sim 1/\sqrt{2}$) error in the re-emission shape from D-scattered photons. Because D scatters are rare at cosmic abundance and the dominant D feature (the absorption dip) depends on opacity rather than re-emission shape, this is acceptable for v1. Per-photon species tracking is also not yet implemented.

The Deuterium Ly α rest wavelength used is $\lambda_D = 1215.337$ (NIST vacuum), ~ 0.33 blueward of $\lambda_H = 1215.668$. In LaRT's velocity convention ($v = -c \Delta\lambda/\lambda_H$, so negative v is blueshifted), the D Ly α line center sits at $v_D \approx -82$ km/s. A reference example with five validation runs and a Jupyter notebook is provided under `examples/lya_HD/` (see `run.sh` and `inspect_HD.ipynb`). One of those runs is configured for direct comparison against Dijkstra, Haiman & Spaans (2006), ApJ 649, 14, Figure 3 ($D/H = 3 \times 10^{-5}$, $\tau \sim 10^6$).

5.3.2 He I 10833 Coherent Triplet Treatment

The He I $\lambda 10830$ multiplet $2^3S_1 \rightarrow 2^3P_{0,1,2}^o \rightarrow 2^3S_1$ is a three-component resonance line (LaRT `line_type=6`). The three upper levels are split by only ~ 0.09 (P_1-P_2 , ~ 2.5 km/s) and ~ 1.25 (P_0-P_2 , ~ 35 km/s), so depending on the gas temperature the components may overlap and interfere.

LaRT offers two treatments of the polarizability parameters (E_1, E_2, E_3) that enter the scattering phase function, selected by the boolean `par%HeI_coherent`:

Toggle	(E_1, E_2, E_3) source
<code>par%HeI_coherent = .false.</code> (default, legacy)	Per-component incoherent constants $(0, 1, 0)$, $(\frac{1}{4}, \frac{3}{4}, \frac{1}{4})$, $(\frac{7}{20}, \frac{13}{20}, \frac{3}{4})$ for P_0, P_1, P_2 from Seon <code>scatter_matrix_v1.1</code> §9.4. Assigned channel-by-channel based on which upper state the photon was sampled into.
<code>par%HeI_coherent = .true.</code>	Frequency-dependent $E_1(\nu), E_3(\nu)$ from the Real- Φ polynomial form of Seon <code>scatter_matrix_HeI</code> memo (eqs. 29 & 30, bottom line), with $E_2 = 1 - E_1$.

The polynomial form is, with $D_m \equiv \nu - \nu_{m1}$ for $m = 0, 1, 2$:

$$E_1(\nu) = \frac{3D_2^2D_0^2 + 7D_0^2D_1^2 + 8D_2D_0D_1^2 + 18D_2D_0^2D_1}{4[D_2^2D_1^2 + 3D_2^2D_0^2 + 5D_0^2D_1^2]}, \quad (1)$$

$$E_3(\nu) = \frac{3D_2^2D_0^2 + 15D_0^2D_1^2 + 8D_2^2D_0D_1 + 10D_2D_0^2D_1}{4[D_2^2D_1^2 + 3D_2^2D_0^2 + 5D_0^2D_1^2]}, \quad (2)$$

$$E_2(\nu) = 1 - E_1(\nu). \quad (3)$$

The common denominator is a positive-definite sum of squares with the three distinct line centers, so the formula is well-defined at every frequency and recovers the v1.1 §9.4 incoherent values exactly at each line center ($E_1, E_3 = 7/20, 3/4$ at P_2 ; $1/4, 1/4$ at P_1 ; $0, 0$ at P_0). It is implemented in `compute_HeI_E_coherent` in `line_mod.f90` (with parallel branches in `scattering_amr.f90` and `line_clump_mod.f90`).

Upper-state selection probability is unchanged. The total cross-section is the incoherent 1:3:5 weighted Lorentzian sum because all $J'-J'$ interference terms cancel in $W_{11} + 2W_{14}$ (the angular integration projects onto the $K=0$ irreducible component). Both modes therefore sample the upper state with the same probability, so the photon-by-photon frequency redistribution is statistically identical. Differences between the two modes appear only in the phase-function asymmetry E_1 (linear Rayleigh polarizability) and the circular polarizability E_3 .

Expected size of the effect. For nebular conditions ($T \sim 10^4$ K) the He-atom Doppler width $\Delta\nu_D \sim 6.4$ km/s *exceeds* the P_1 – P_2 splitting (~ 2.5 km/s), so those two components are effectively merged and the inter-line interference is small in the line core. The escape spectrum is essentially unchanged between the two modes. Polarization differences are visible mainly at high τ ($\tau \gtrsim 100$) and in the line wings where photons sample the inter-line frequency range; integrated polarization profiles differ by ~ 5 – 15% (relative) at $\tau \sim 10^3$. The effect would be larger at lower temperature ($T \lesssim 10^3$ K), where the P_1 and P_2 components are resolved.

A reference test suite is provided under `examples/HeI_coherent_test/` with 20 runs covering $\{\text{point, uniform-sphere source}\} \times \{\tau = 0.1, 1, 10, 100, 1000\} \times \{.false., .true.\}$ and a comparison notebook `compare.ipynb`.

5.4 Frequency and Wavelength Grid

Internally, photon frequencies are tracked in the dimensionless Doppler variable:

$$x = \frac{\nu - \nu_0}{\Delta\nu_D}, \quad \Delta\nu_D = \frac{\nu_0}{c} b_D, \quad b_D = \sqrt{\frac{2kT}{m} + b_{\text{turb}}^2}, \quad (4)$$

where ν_0 is the line-center frequency, T is the local gas temperature, m is the atomic/ionic mass, and $b_{\text{turb}} = \text{par\%bturb}$ is a spatially uniform turbulent velocity added *in quadrature* to the local thermal width ($\text{par\%bturb} \leq 0$, the default, recovers the pure thermal width). The same b_D , evaluated with the *local* cell temperature, sets the line opacity, the Voigt parameter, the scattering redistribution, and the velocity-to-Doppler conversion identically on the Cartesian, AMR, and clump grids. Output spectra are binned onto a uniform grid in x by default. A wavelength or velocity grid can be requested instead:

Parameter	Default	Description
par%nxfreq	121	Number of spectral bins
par%xfreq_min, par%xfreq_max	auto	Limits of x -grid
par%xfreq0	0.0	Frequency offset applied to all emitted photons
par%nwavelength	0	Number of wavelength bins (overrides x -grid)
par%wavelength_min, par%wavelength_max	—	[Å]
par%nvelocity	0	Number of velocity bins (overrides x -grid)
par%velocity_min, par%velocity_max	—	[km s ⁻¹]

The range is set automatically from the mean medium temperature when no limits are provided.

6. Medium Properties

6.1 Optical Depth Normalization

The overall opacity is scaled by one of the following (specify at most one):

Parameter	Description
par%taumax	Line-center τ_0 from the grid center to the $+z$ boundary along the z -axis
par%tauhomo	Same, assuming a uniform density field (scales an inhomogeneous field proportionally)
par%N_HI_max	H I column density from center to $+z$ boundary [cm ⁻²]
par%N_HI_homo	Same, assuming uniform density

For AMR grids, $\text{par\%taumax}$ is the pole-to-face optical depth (box center $\rightarrow +z$ face along z), consistent with the Cartesian convention for a sphere of radius $R = L/2$.

6.2 Temperature

Parameter	Default	Description
par%temperature	10 ⁴ K	Uniform gas temperature
par%temperature0	$= T$	Temperature for <code>spectral_type='voigt0'</code> source
par%bturb	0.0	Turbulent velocity dispersion added in quadrature to thermal width [km s ⁻¹]
par%temp_file	''	3-D temperature cube [K]; FITS or HDF5

6.3 Density

Analytic profiles. Without external files, the density is set analytically via the opacity scale parameters (par%taumax etc.) combined with an optional radial profile:

Profile type	Key parameter	Functional form
Uniform (sphere or box)	par%rmax	Constant n_{H} inside $r < r_{\text{max}}$
Spherical shell	par%rmin, par%rmax	Constant n_{H} between r_{min} and r_{max}
Exponential	par%density_rscale	$n \propto \exp(-r/r_s)$
Power law	par%density_alpha	$n(r) \propto (r_{\text{max}}/r)^{\alpha_D}$; $\alpha_D=0$: uniform, $\alpha_D=2$: isothermal

External gridded data (FITS or HDF5).

The reader (read_grid_data.f90) detects the format from the filename extension (.fits[.gz], .h5, or .hdf5) and uses the same code path for both. HDF5 inputs expect the array at /<section>/data (any section name); shape information is read from the dataset itself, so FITS-style NAXIS keywords are not required.

Parameter	Description
par%dens_file	3-D H I number density cube [cm^{-3}]
par%temp_file	3-D temperature cube [K]
par%velo_file	4-D velocity cube (v_x, v_y, v_z) [km s^{-1}]
par%emiss_file	3-D emissivity cube [$\text{photons s}^{-1} \text{cm}^{-3}$]
par%input_field	Base name: reads <base>.dens.fits.gz, .temp., .velo. (legacy convention; FITS-only)

If par%taumax (or par%N_HI_max) is specified alongside a density file, the density is rescaled so that the target optical depth is achieved.

6.4 Velocity Fields

par%velocity_type selects the bulk velocity model. The default (empty string) is a static medium.

Value	Key parameters	Velocity law
'hubble'	par%Vexp	$\mathbf{v}(r) = V_{\text{exp}}(r/r_{\text{max}})\hat{r}$; $V_{\text{exp}} > 0$: outflow
'constant_radial'	par%Vexp	$\mathbf{v}(r) = V_{\text{exp}}\hat{r}$ (constant magnitude)
'power_law'	par%Vexp, par%velocity_alpha	$\mathbf{v}(r) = V_{\text{exp}}(r/r_{\text{max}})^{\alpha_V}\hat{r}$
'linear_decelerate'	par%Vexp	$\mathbf{v}(r) = V_{\text{exp}} \frac{r_{\text{max}} - r}{r_{\text{max}} - r_{\text{min}}} \hat{r}$
'parallel_velocity'	par%Vx, par%Vy, par%Vz	Uniform $\mathbf{v} = (V_x, V_y, V_z)$ [km s ⁻¹]
'ssh'	par%Vpeak, par%rpeak, par%DeltaV	Bilinear profile (Song, Seon & Hwang 2020)
'rotating_solid_body'	par%Vrot, par%rinner	Solid-body rotation
'rotating_galaxy_halo'	par%Vrot, par%q	Galaxy halo rotation curve

The 'ssh' profile rises linearly to V_{peak} at r_{peak} , then shifts by ΔV toward r_{max} :

$$v(r) = \begin{cases} V_{\text{peak}} r/r_{\text{peak}} & r < r_{\text{peak}} \\ \left[V_{\text{peak}} + \Delta V \frac{r - r_{\text{peak}}}{r_{\text{max}} - r_{\text{peak}}} \right] \hat{r} & r \geq r_{\text{peak}} \end{cases} \quad (5)$$

A velocity FITS file (par%velo_file) may be combined with a systematic par%velocity_type.

Other velocity-related parameters:

Parameter	Default	Description
par%Omega	28.0	Shear rate [(km s ⁻¹) kpc ⁻¹] for TIGRESS data
par%recoil	.false.	Include photon-recoil frequency shift in scattering
par%core_skip	.false.	Enable resonance-core skipping acceleration
par%core_skip_global	.false.	With core_skip: use the old volume-averaged x_{crit} instead of the per-cell formula (benchmark)

Core-skip details. When `core_skip = .true.` and `core_skip_global = .false.` (the default combination), LaRT recomputes the core-skip threshold x_{crit} at every scattering using the local cell:

$$a\tau_{\text{cell}} = a_V(\text{cell}) \rho_{\text{hokap}}(\text{cell}) \Delta\ell_{\text{face}}, \quad x_{\text{crit}} = \begin{cases} (a\tau_{\text{cell}})^{1/3}/5 & a\tau_{\text{cell}} > 1 \\ 0 & \text{otherwise} \end{cases}$$

where $\Delta\ell_{\text{face}}$ is the photon's distance to the nearest cell boundary. This is Smith et al. (2015, Eq. 35) applied per cell, following the RASCAS convention. Inhomogeneous boxes such as a dense galactic core in a sparse hot halo require this per-cell treatment; the `core_skip_global = .true.` branch uses a single volume-averaged x_{crit} that collapses to zero whenever $\langle a\tau \rangle < 1$ and is retained only for comparison with the older LaRT behavior.

7. AMR Grid Mode

7.1 Enabling AMR Mode

v2.00:

```
1 par%use_amr_grid = .true.
2 par%amr_file      = 'grid.h5'
```

v2.10 (preferred):

```
1 par%grid_type = 'amr'
2 par%amr_file  = 'grid.h5'
```

LaRT reads only the *generic AMR format* (text, FITS, or HDF5). Simulation data from RAMSES or other codes must first be converted using the standalone converter (`convert_ramses_to_generic.x` or the Python equivalent). The parameter `par%amr_type` defaults to 'generic'; setting it to 'ramses' is no longer supported and produces an error directing the user to the converter.

Parameter	Default	Description
<code>par%amr_type</code>	'generic'	Only 'generic' is supported
<code>par%amr_file</code>	''	Path to generic AMR file (.h5, .fits.gz, .dat)

7.2 Generic AMR Data Format

The generic AMR format stores leaf-cell data in text, FITS binary table, or HDF5. Nine columns are mandatory; six optional columns provide pre-computed physics quantities that LaRT can use directly.

Mandatory columns (9):

Column	Unit	Description
x, y, z	code units	Leaf cell center coordinates
level	—	AMR refinement level
nH (or gasDen, dens)	cm ⁻³	Total hydrogen number density
T	K	Gas temperature
vx, vy, vz	km s ⁻¹	Gas velocity

Optional columns (detected by name in HDF5/FITS):

Column	Unit	Description
metallicity	mass fraction	Raw Z from simulation
xHI	dimensionless	Neutral hydrogen fraction
n_e	cm ⁻³	Electron density
n_ion	cm ⁻³	Scattering ion density (metal lines)
emissivity	cm ⁻³ s ⁻¹	Line emissivity rate (Ly α , C IV, or O VI)
ndust	cm ⁻³	Dust pseudo-number density

When an optional column is present in the file, LaRT uses it directly. When absent, LaRT computes the quantity internally using the model selected by the extended physics parameters (see below).

Plain text format: a header followed by one line per leaf cell.

```

1 # comment lines start with #
2 BOXLEN  100.0
3 NLEAF   8192
4 # x   y   z   level  nH   T   vx   vy   vz   metallicity   xHI
5 50.5  50.5  50.5  1   1.0e-3  1.0e4  0.0  0.0  0.0  0.01  0.95
6 ...

```

Leaf-cell coordinates are given with the box origin at the corner, $[0, \text{BOXLEN}]$. The reader internally re-centers to $[-L/2, +L/2]$; *do not* pre-shift the data.

7.3 Extended AMR Physics Parameters

These parameters control how LaRT computes physical quantities when the corresponding optional column is absent from the generic file. All defaults reproduce the existing behavior.

Parameter	Default	Description
par%ionization_model	'cie_formula'	Neutral fraction model: 'cie_formula' (single-formula CIE), 'cie_table' (Gnat & Sternberg 2007), 'full_neutral' ($x_{\text{HI}} = 1$), 'from_file' (require xHI column)
par%dust_model	'global_dgr'	Dust model: 'global_dgr' (use par%DGR), 'laursen09' (metallicity-based, Laursen et al. 2009), 'from_file' (require ndust column)
par%emissivity_model	'none'	Ly α emissivity: 'none', 'caseB' (Case B recombination + collisional excitation), 'from_file'
par%ion_model	'none'	Metal-line ion density: 'none', 'solar_cie' (solar abundances $\times$ CIE ion fractions), 'from_file'
par%metallicity_global	-1	Global metallicity (Z , mass fraction). Used by 'laursen09' and 'solar_cie' when no metallicity column is present.
par%Z_ref	0.0134	Reference solar metallicity (Asplund et al. 2009)
par%f_ion_dust	0.01	Ionized-gas dust fraction (Laursen et al. 2009)

7.4 Tau Normalization Convention

For AMR, par%taumax is the line-center optical depth from the box center to the $+z$ face along the z -axis. This matches the Cartesian sphere convention where par%taumax is the center-to-surface optical depth, provided $r_{\text{max}} = L/2$.

7.5 Converting RAMSES Data

RAMSES binary snapshots must be converted to the generic format before use with LaRT. Two converters are provided:

Fortran:

```
1 make convert_ramses
2 convert_ramses_to_generic.x /path/to/ramses 91 output.h5 kpc
3 convert_ramses_to_generic.x /path/to/ramses 91 output.h5 kpc \
4   --compute-physics --metallicity 0.0134
```

Python:

```
1 python convert_ramses_to_generic.py /path/to/ramses 91 -o output.h5 \
2   --compute-physics --metallicity 0.0134
```

The `--compute-physics` flag pre-computes $xH\text{I}$, n_e , and emissivity using the CIE formula. If `--metallicity` is given, n_{dust} is also computed via the Laursen et al. (2009) model.

7.6 Converting Illustris/IllustrisTNG Data

Illustris and IllustrisTNG simulations use the AREPO moving-mesh code, which stores gas data on an unstructured Voronoi tessellation. Since LaRT requires an octree AMR grid, the Voronoi data must be resampled onto an adaptive octree. The converter `convert_illustris_to_generic.py` handles this automatically:

```
1 python convert_illustris_to_generic.py cutout.hdf5 -o galaxy.h5 \
2   --output-unit kpc --level-max 8 --compute-physics \
3   --ionization from_illustris --boxsize 200
```

The converter reads `PartType0` fields from an HDF5 snapshot or API cutout, converts Illustris comoving+ h code units to physical (kpc, cm^{-3} , K, km s^{-1}), builds an adaptive octree via density and velocity gradient refinement, and assigns leaf-cell properties via nearest-neighbor KD-tree lookup (physically exact for Voronoi tessellation).

Key options:

Flag	Description
--grid-type	Output grid: amr (adaptive octree, default) or cartesian (uniform N^3)
--ngrid	Cells per side for --grid-type cartesian (default 128)
--level-max	Maximum octree level (default 8; AMR mode only)
--dens-threshold	Density gradient threshold (default 0.3; AMR mode only)
--match-resolution	Also refine to match local Voronoi cell size (AMR mode only)
--compute-physics	Compute xHI, n_e, emissivity, ndust
--emissivity-line	lya (Case B, default), civ (CIE CIV), ovi (CIE OVI)
--ionization	from_illustris (TNG native), cie, or full_neutral
--sfr-treatment	cap-temperature (default), exclude, or as-is
--api-key	Download cutout from TNG public API
--plot	Diagnostic slice plot

The --sfr-treatment cap-temperature option (default) caps the temperature of star-forming cells at 2×10^4 K, since their InternalEnergy reflects an effective equation of state, not the actual ISM temperature.

Cutouts can be downloaded directly from the TNG public API (tng-project.org; free registration required):

```
1 python convert_illustris_to_generic.py \
2     --api-key YOUR_KEY --simulation TNG50-1 --snap 99 \
3     --subhalo-id 0 -o galaxy.h5 --compute-physics
```

When the converted grid file contains an emissivity column, `source_geometry = 'diffuse_emissivity'` will automatically use it as the photon source distribution without requiring `par%emiss_file` to be set. Alternatively, set `par%emissivity_model = 'caseB'` to compute the emissivity on the fly from n_{H} , T , and x_{HI} .

Cartesian grid output. In addition to the adaptive octree, the converter can produce a uniform $N \times N \times N$ Cartesian grid:

```
1 python convert_illustris_to_generic.py cutout.hdf5 \
2     -o galaxy_cart.h5 --grid-type cartesian --ngrid 128 \
3     --output-unit kpc --compute-physics --ionization cie
```

The output file stores all quantities as 3D image sections (one per physical variable), read by LaRT via `par%cart_file`:

```
1 par%cart_file      = 'galaxy_cart.h5'
2 par%distance_unit = 'kpc'
```

This uses the standard Cartesian raytrace with no AMR overhead. Optional columns (xHI, emissivity, ndust) are applied automatically when present.

For full details of the Illustris data format, unit conversions, and the Voronoi-to-octree algorithm, see [docs/Illustris_data_structur](#)

7.7 Source Geometries in AMR Mode

In v2.00, AMR supports 'point' and 'uniform'. In v2.10, most Cartesian geometries (Gaussian, exponential, Sérsic, star_file, etc.) are also available; refer to §5.1 and the v2.10 release notes.

8. Clumpy Medium Mode

8.1 Overview

Clumpy medium mode places N spherical clumps of radius r_c inside a spherical outer domain of radius R . The inter-clump medium is vacuum. All opacity resides inside the clumps. By default, clump positions are drawn by Random Sequential Adsorption (RSA) so that clumps never overlap. Optionally, overlapping clumps can be generated or loaded from an external FITS file; LaRT detects overlaps automatically after placement and engages an event-based multi-component raytrace in that case (see §8.6).

8.2 Enabling Clump Mode

v2.00:

```
1 par%use_clump_medium = .true.
```

v2.10 (preferred):

```
1 par%grid_type = 'clump'
```

8.3 Clump Parameters

Parameter	Default	Description
par%rmax	1.0	Outer sphere radius (code units)
par%rmin	-999	Inner cavity radius (the default sentinel is treated as 0). Clumps are placed only in the shell $r_{\min} \leq r \leq r_{\max}$
par%clump_radius	—	Individual clump radius (same units)
par%clump_tau0	—	Line-center τ_0 from clump center to its surface
par%clump_f_cov	—	Covering factor $C_f = Nr_c^2/R^2$
par%clump_f_vol	—	Volume filling factor $f_V = N(r_c/R)^3$
par%clump_N_clumps	—	Total number of clumps N
par%temperature	10^4 K	Clump gas temperature
par%clump_sigma_v	0.0	Gaussian σ of random bulk clump velocities [km s^{-1}]
par%save_clump_info	.false.	Write clump positions and velocities to a FITS binary table
par%clump_input_file	''	If non-empty, load the clump population from this FITS file (see §8.5) instead of generating it internally
par%cone_opening	0.0	Biconical outflow half-opening angle [degrees]. When > 0 , clumps are placed only inside the bicone (both hemispheres). With par%clump_fully_inside, clumps are additionally rejected if they protrude outside the cone boundary.
par%clump_fully_inside	.true.	If .true. (default), reject any candidate whose outermost edge would protrude past the outer sphere ($d_{\text{center}} + r_c > R$) or into the inner cavity ($d_{\text{center}} - r_c < r_{\min}$); if .false., restore the legacy behavior in which only the center is required to lie in $[r_{\min}, R]$
par%clump_allow_overlap	.false.	If .true., skip RSA overlap rejection during internal generation: clumps are placed at uniformly random positions. After placement check_has_overlap() is called automatically; if any overlapping pair is found the overlap-aware raytrace is engaged. See §8.6

Specify exactly one of par%clump_f_cov, par%clump_f_vol, or par%clump_N_clumps.

For the per-clump opacity, specify *one* of par%clump_tau0, par%clump_NHI, or par%clump_nH. As a convenience, when none of those three is set the system-level radial-sightline targets par%taumax and par%N_HI_max are also accepted for procedurally generated clumps: LaRT back-solves the peak per-clump opacity so the realised radial sightline from the box centre to $r = R$ matches the requested par%taumax (or par%N_HI_max). In the uniform-clump case the back-solve uses the closed-form $\tau_{\max} = \frac{4}{3} f_{\text{cov}} \tau_{c,0}$ (and similarly for $N_{\text{HI,max}}$); in profile mode it uses the radial quadra-

ture derived from `par%clump_radius_profile`, `par%clump_density_profile`, and `par%clump_number_profile`. `par%N_gasmax` is treated as an alias for `par%N_HImax` in clump mode (HI-only line). When a pre-built clump population is loaded via `par%clump_input_file`, any one of `par%taumax`, `par%tau homo`, `par%N_HImax` (= `par%N_gasmax`), or `par%N_HI homo` (= `par%N_gashomo`) may be supplied as the rescale target. Every `cl_rho kap(i)` is multiplied by a single factor α so the chosen target is exactly reproduced; the other three scalars are derived consistently from the rescaled distribution. Priority order when several targets are given: `par%taumax` > `par%tau homo` > `par%N_gasmax` > `par%N_gashomo`. See the *Clumpy Medium Algorithm Description*, §“System-level fallback”.

When `par%rmin` is set, `par%clump_f_vol` and `par%clump_f_cov` are interpreted relative to the shell volume $V_{\text{shell}} = \frac{4}{3}\pi(R^3 - r_{\text{min}}^3)$ and the shell sightline factor $R^2 + Rr_{\text{min}} + r_{\text{min}}^2$. With `par%rmin` unset (default -999, treated as 0) these reduce exactly to the original $(R/r_c)^3$ and $(R/r_c)^2$ formulas, preserving prior runs. When `par%clump_fully_inside` is on, the realized filling factors fall slightly below the requested target whenever r_c is a non-trivial fraction of $R - r_{\text{min}}$, since the available placement volume shrinks from $\frac{4}{3}\pi(R^3 - r_{\text{min}}^3)$ to $\frac{4}{3}\pi[(R - r_c)^3 - (r_{\text{min}} + r_c)^3]$. Specify a slightly larger target if exact matching is required. See the *Clumpy Medium Algorithm Description* document, “Fully-Inside Constraint” subsection, for the placement details.

8.4 Spatially-varying clump properties

Three optional radial-profile inputs let the per-clump radius, internal HI density, and spatial number density vary with distance from the box center. Built-in shapes are ‘constant’ (default), ‘powerlaw’, ‘gaussian’, ‘exponential’, and ‘file’ (5-column ASCII table). Setting any flag to a value other than ‘constant’ switches `par%init_clumps` to the profile-aware path.

Parameter	Default	Description
<code>par%clump_radius_profile</code>	‘constant’	shape of $r_c(r)$
<code>par%clump_radius_alpha</code>	0.0	power-law index for $r_c(r)$
<code>par%clump_radius_r0</code>	0.0	power-law scale for $r_c(r)$; $\leq 0 \Rightarrow r_{\text{max}}$
<code>par%clump_density_profile</code>	‘constant’	shape of internal $n_H(r)$
<code>par%clump_density_alpha</code>	0.0	power-law index for $n_H(r)$
<code>par%clump_density_r0</code>	0.0	power-law scale for $n_H(r)$; $\leq 0 \Rightarrow r_{\text{max}}$
<code>par%clump_number_profile</code>	‘constant’	shape of spatial $n_{cl}(r)$
<code>par%clump_number_alpha</code>	0.0	power-law index for $n_{cl}(r)$
<code>par%clump_number_r0</code>	0.0	power-law scale for $n_{cl}(r)$; $\leq 0 \Rightarrow r_{\text{max}}$
<code>par%clump_profile_file</code>	‘’	5-column tabulated profile (used when any axis = ‘file’)

For non-uniform profiles the closed-form $f_{\text{cov}} / f_{\text{vol}} / N$ relation does not hold. Instead the profile-aware quadrature back-solves the spatial normalization so that whichever of `par%clump_f_cov`, `par%clump_f_vol`, or `par%clump_N_clumps` is provided is matched, and the other two are reported on stdout. See the *Clumpy Medium Algorithm Description* document for the integral definitions and the implementation outline.

8.5 Standalone clump generator (make_clumps.x)

A self-contained Fortran driver, `make_clumps.x`, builds a clump population without running the photon transport. It is no longer linked against MPI or the LaRT physics stack; the only run-time dependency is the FITS/HDF5 output facade. All parameters are passed as command-line flags (one flag per `par%clump_*` field plus the relevant global parameters):

```

1 make make_clumps                                # builds make_clumps.x
2 ./make_clumps.x --rmax 1.0 --clump_radius 1e-3 \
3                 --clump_f_cov 5 --clump_tau0 1e3 \
4                 --temperature 1e4 --line_id ly_alpha \
5                 --seed 12345 -o run_clumps.fits.gz

```

Run `./make_clumps.x --help` for the full flag list. The output filename's extension chooses the container: `.fits` / `.fits.gz` produce a FITS binary table, `.h5` / `.hdf5` produce an HDF5 table. In both cases the header keywords and column schema match `$??` so the file is directly consumable by LaRT through `par%clump_input_file`.

A pure-Python equivalent, `python/make_clumps.py`, ships in the same repository; the two drivers expose essentially the same flag set and write the same schema. All sampling paths are supported in both: uniform RSA, radial profiles (powerlaw, gaussian, exponential), biconical geometry (`--cone_opening`, including the `clump_fully_inside` angular-subtension test), every opacity-input mode (`--clump_tau0` / `--clump_NHI` / `--clump_nH` / `--taumax` / `--N_HI_max`), and the same systematic velocity families (hubble, constant_radial, power_law, linear_decelerate, parallel_velocity, ssh, rotating_solid_body, rotating_galaxy_halo) on top of the Gaussian `par%clump_sigma_v` component.

```

1 python python/make_clumps.py --rmax 1.0 --clump_radius 1e-3 \
2     --clump_f_cov 5 --clump_tau0 1e3 --temperature 1e4 \
3     --line_id ly_alpha --seed 12345 -o run_clumps.fits.gz

```

Both drivers are deterministic in `--seed`: re-running with the same flag set and the same seed reproduces the table bit-for-bit. The two drivers use independent RNG implementations, however, so for a given seed the Fortran and Python outputs are statistically equivalent but not byte-identical; choose whichever fits the host environment. Either output is dropped into a LaRT run by setting `par%clump_input_file`.

To consume the file, set `par%clump_input_file` on the LaRT input:

```

1 &parameters
2   par%grid_type      = 'clump'
3   par%rmax           = 1.0
4   par%clump_radius   = 0.01
5   par%clump_tau0     = 1.0e3
6   par%temperature    = 1.0e4
7   par%no_photons     = 1e6
8   par%clump_input_file = 'clumps_clumps.fits.gz'
9 /

```

When `par%clump_input_file` is non-empty, `init_clumps` skips the profile-detection and RSA-generation paths entirely and reads the population from the file. `read_clumps_fits` accepts files with any subset of the per-clump columns `R_CLUMP/RHOKAP/TEMP`; missing columns are filled from the corresponding header keyword (`CL_RAD/RHOKAP/TEMP_CL`). For the RHOKAP slot the column names `DENSITY` and `DENS` are also accepted as aliases, in priority order `RHOKAP` → `DENSITY` → `DENS`. The interpretation depends on which name was used: `RHOKAP` (and the header-keyword fallback) is taken as line-centre opacity per code unit and stored directly into κ_i , while `DENSITY` or `DENS` are taken as the neutral-hydrogen number density $n_{HI,i}$ [cm^{-3}] and converted to opacity via $\kappa_i = n_{HI,i} \cdot \text{line\%cross0} / \Delta\nu_{D,i} \cdot \text{par\%distance2cm}$. The conversion requires `par%distance_unit` (and therefore `par%distance2cm`) to be set; if not, the read path emits a warning and falls back to treating the values as opacity.

The clump file also carries `DISTUNIT` (`par%distance_unit`) and `DIST_CM` (`par%distance2cm`) keywords written by LaRT itself. On read, these are adopted into `par%distance_unit` / `par%distance2cm` when the user did NOT supply them in the input namelist; the criterion for “user-supplied” is `par%distance_unit != ''` and `par%distance2cm > 1.0`, so the file values take effect only when the namelist leaves both at their defaults. This lets a clump file generated with `par%distance_unit = 'kpc'` be replayed in a downstream run without having to repeat that setting, and gives a meaningful `par%distance2cm` to the `DENSITY/DENS` opacity conversion above. External pipelines that write LaRT-compatible HDF5 clump files must store `DISTUNIT` as a fixed-length ASCII string (`h5py: string_dtype(encoding='ascii', length=N)`); variable-length strings are detected on read and silently skipped (a warning is printed).

If any of `par%taumax`, `par%tauhomo`, `par%N_HI_max` (`=par%N_gasmax`), or `par%N_HI_homo` (`=par%N_gashomo`) is specified together with `par%clump_input_file`, every loaded `cl_rhokap(i)` is multiplied by a single factor $\alpha = \text{target}/\text{realised}$ so the chosen system-level scalar matches the input target; the other three scalars are then derived consistently from the rescaled population. Priority among multiple targets: `par%taumax > par%tauhomo > par%N_gasmax > par%N_gashomo`. If none of the four is set, the loaded opacities are taken as-is and the four scalars are reported as computed from the file. See *Clumpy Medium Algorithm Description*, §“System-level fallback” and §“Derived Global Optical-Depth Quantities” for the formulas.

8.6 Overlapping Clumps

By default LaRT guarantees non-overlapping clumps through RSA placement and all raytrace paths assume this property. Two situations can produce overlapping populations:

1. A clump population loaded from an external FITS file via `par%clump_input_file` may contain overlapping clumps.
2. Setting `par%clump_allow_overlap = .true.` skips the RSA rejection step so that internally generated clumps are placed at uniformly random positions and will typically overlap whenever the clumps are large relative to the domain.

After the CSR acceleration grid is built, `check_has_overlap()` scans all clump pairs. If any overlapping pair is found it prints

```
1 Overlap check: overlapping clumps detected -- using overlap-aware raytrace.
```

sets an internal flag `has_overlap = .true.`, and routes all subsequent raytrace calls to the overlap-aware variants automatically. No input change is required; the detection and dispatch happen at run time.

Physical model for overlapping clumps

Two clumps sharing a volume with different bulk velocities $\mathbf{v}_1$ and $\mathbf{v}_2$ are treated as two independent gas components coexisting at the same spatial location. The total line opacity at global-frame frequency x is

$$\kappa_{\text{total}}(x) = \sum_i \rho \kappa_i H(x - u_{\text{los},i}, a_i), \quad u_{\text{los},i} = \frac{\mathbf{v}_i \cdot \hat{\mathbf{k}}}{\Delta \nu_{D,i}}. \quad (6)$$

When a scattering event occurs in an overlap region the interacting atom belongs to one clump, chosen by opacity-weighted sampling $P(\text{clump } i) = \kappa_i / \kappa_{\text{total}}$.

Quick-start example

Large clumps with low covering factor guarantee many overlaps while keeping the run fast:

```
1 par%grid_type           = 'clump'
2 par%clump_radius        = 0.3
3 par%clump_f_cov         = 2.0
4 par%clump_allow_overlap = .true.
5 par%no_photons          = 1e5
6 par%spectral_type       = 'voigt'
```

Reference inputs: `clump_overlap_mono.in`, `clump_overlap_sigv.in`, `clump_overlap_peel.in` in `examples/clump_sphere/`.

For a full algorithmic description see the *Clumpy Medium Algorithm Description* document, Section 15 “Overlap-Aware Raytrace”.

8.7 Clump Velocity

A systematic flow can be added on top of the random component (`par%clump_sigma_v`) via `par%velocity_type`. Supported values in clump mode: `'hubble'`, `'constant_radial'`, `'power_law'`, `'linear_decelerate'`, `'parallel_velocity'`, `'ssh'`, `'rotating_solid_body'`, `'rotating_galaxy_halo'`.

9. Observers and Peel-off

The peel-off (peeling-off) technique shoots a photon toward the observer at every scattering event, allowing spectral images to be constructed without noise. It is activated by setting `par%nxim` and `par%nyim` to non-zero values.

9.1 Observer Placement

Direction vector:

```

1 par%obsx = 0.0, par%obsy = 0.0, par%obsz = 1.0  ! look down +z axis
2 par%distance = 100.0                               ! observer-center separation

```

Spherical angles:

```

1 par%alpha = 45.0  ! azimuthal angle [deg]
2 par%beta  = 30.0  ! elevation angle [deg]
3 par%distance = 100.0

```

par%gamma sets the rotation of the detector plane about the line of sight. If par%distance is omitted it defaults to $100 \times r_{\max}$ (or the maximum box half-extent).

9.2 Multiple Observers

Multiple observers are specified as arrays:

```

1 par%alpha = 0.0  45.0  90.0  135.0
2 par%beta  = 0.0   0.0   0.0   0.0

```

Up to 99 simultaneous observers are supported.

9.3 Image Parameters

Parameter	Default	Description
par%nxim, par%nyim	0	Image pixel counts; peel-off enabled when > 0
par%dxim, par%dyim	auto	Pixel angular size [deg] (defaults to cover the full grid)

When par%dxim/par%dyim are not set, the pixel size is chosen to cover the full sphere:

$$\delta\theta = \frac{1}{N_x/2} \arcsin\left(\frac{r_{\max}}{d}\right). \quad (7)$$

9.4 Peel-off Output Flags

Parameter	Default	Description
par%save_peeloff	.false.	Save 1-D peel-off spectra per observer
par%save_peeloff_2D	.false.	Save peel-off spectral image cubes (x - y - ν)
par%save_peeloff_3D	.false.	Save full 3-D peel-off data
par%save_sightline_tau	.false.	Compute deterministic sight-line maps from each observer pixel: gas $\tau(\nu)$, N_{HI} , and dust τ (when par%DGR > 0). See §9.5.

9.5 Sight-line Optical-depth and Column-density Maps

When `par%save_sightline_tau = .true.`, LaRT integrates the opacity along each observer pixel's line of sight and writes a separate FITS file `<base>_tau<obs>.fits.gz` for every observer. The file contains:

HDU	Content
TAU_gas	gas optical depth $\tau(\nu, x, y)$
N_gas	gas column density $N_{\text{HI}}(x, y)$
TAU_dust	dust optical depth $\tau_{\text{dust}}(x, y)$ (only when <code>par%DGR > 0</code>)

The pixel grid uses the same WCS as the peel-off image (`par%nxim`, `par%nyim`, `par%dxim`, `par%dyim`); the frequency axis matches the spectrum output (`par%nxfreq`, `par%velocity_min`, `par%velocity_max`). All three grid modes (Cartesian, AMR, clump) are supported with the appropriate raytrace dispatch.

Standalone driver `make_sightline_tau.x`

The same sight-line integration is exposed as a standalone executable that skips the photon-transport stage entirely:

```
1 make sightline_tau                # builds make_sightline_tau.x
2 mpirun -np N ./make_sightline_tau.x sl.in    # writes <base>_tau.fits.gz
```

`sl.in` is an ordinary LaRT input file; the driver internally sets `par%save_peeloff` and the sight-line flag to `.true.` and otherwise honors every namelist parameter. Useful for parameter sweeps where only the deterministic maps are needed (no Monte-Carlo transport). Unlike `make_clumps.x`, multiple MPI ranks do speed this driver up: pixels are divided across `par%nproc` via `loop_divide` and reduced with `reduce_mem`.

The folder `examples/sightline_tau/` contains a reference set of six runs that exercise every grid mode and observer configuration the driver supports:

File	Grid mode	Observer
<code>car_sightline.in</code>	Cartesian	outside (rectangular FITS)
<code>amr_sightline.in</code>	AMR	outside (rectangular FITS)
<code>clump_sightline.in</code>	clumpy sphere	outside (rectangular FITS)
<code>car_inside.in</code>	Cartesian	inside (HEALPix all-sky)
<code>amr_inside.in</code>	AMR	inside (HEALPix all-sky)

The shipped `run.sh` executes all six in sequence; `make_amr_sphere_data.py` regenerates the AMR data file from a uniform-sphere physics model. The companion notebook `plot_sightline_tau.ipynb` displays the column-density maps, selected $\tau(x, y, \nu)$ slices, the all-sky HEALPix mollview maps and the sky-mean line profile.

tau_huge cap (sight-line vs. photon transport). The AMR raytrace routines used by the photon-transport stage early-exit when the running optical depth exceeds the double-precision $\exp(-\tau)$ underflow point ($\tau_{\text{huge}} \approx 745.2$); accumulating further contributes nothing to the photon's transmission. The sight-line variant

raytrace_to_edge_tau_gas_amr exists precisely to *report* the integrated τ , however large, and therefore runs without that cap (matching the Cartesian raytrace_to_edge_tau_gas in raytrace_car.f90). The same applies to raytrace_to_dist_tau_gas_amr used by the inside-observer HEALPix path.

9.6 Inside Observer (HEALPix All-Sky)

Setting `par%inside > 0` switches the entire output and peel-off chain to the *inside-observer* path: the observer is placed at a chosen position (o_x, o_y, o_z) *inside* the medium, and each escaping photon is binned into an all-sky HEALPix map by the direction from the observer back to the photon. This mode is intended for “galactic” observations – e.g., placing the observer at the Sun’s position inside a model of the Galactic ISM, or inside a simulated galaxy.

Activation:

```

1 par%inside = 8                ! HEALPix Nside parameter; total npix = 12*inside^2
2 par%obsx   = 0.3              ! observer position (code units)
3 par%obsy   = 0.0
4 par%obsz   = 0.0
5 par%save_peeloff = .true.
6 par%save_peeloff_2D = .true. ! also save the 2D map (sum over frequency)

```

When `par%inside > 0`, `setup.f90` automatically sets `par%observer_located_inside = .true.` and switches:

- the observer arrays from rectangular images to HEALPix all-sky maps with $N_{\text{pix}} = 12N_{\text{side}}^2$;
- the peel-off routines to the `_inside` variants (which use `vec2pix` to bin the peeled weight by observer-pointing direction);
- `output_reduce`, `output_normalize` and `write_output` to the `_inside` variants.

In v2.10 the dispatch is performed by the type-bound methods of `grid_base_type` (`grid_class.f90`); the `grid_car_type` and `grid_amr_type` concrete types each branch on `par%observer_located_inside` and forward to the appropriate `_inside` or `_outside` concrete routine.

Multiple observers. As with the rectangular path, up to 99 observer positions are supported by passing arrays:

```

1 par%inside = 8
2 par%obsx   = 0.0  0.3  0.0
3 par%obsy   = 0.0  0.0  0.5
4 par%obsz   = 0.0  0.0  0.0

```

Output files. In inside-observer mode the `_obs.fits.gz` file holds 2D HEALPix arrays of shape `(npix, nxfreq)` with extensions `Direct` (unscattered direct contribution) and `Scattered` (scattered contribution). When `par%save_peeloff_2D = .true.` a separate `_obs2D.fits.gz` file is written with shape `(npix,)` holding the same arrays summed over frequency. The HEALPix maps use the RING ordering convention. The Spectrum HDU (`Jout`, `Jin`, `Jabs`) is always written regardless of observer mode.

Source scope. The inside-observer path supports the same internal sources as the standard path – point, uniform, gaussian, exponential, sersic, diffuse_emissivity – but *not* the external-illumination geometries (point_illumination, stellar_illumination, plane_illumination), which are only physically meaningful for an observer outside the medium. Stokes polarization is not yet implemented for the inside-observer path.

Grid-mode support. The HEALPix inside-observer path is available for both Cartesian (par%grid_type = 'car') and AMR (par%grid_type = 'amr') grids. In AMR mode, dedicated _inside_amr peeling, output and sight-line routines are used (see the *LaRT AMR Description* document for implementation details). Clumpy medium (par%grid_type = 'clump') is not currently supported with the inside-observer path.

Sightline-tau output. When the inside-observer path is active, an additional _sightline.fits.gz file is produced that records τ and $N(\text{HI})$ from the observer position to each of the six box faces ($\pm x, \pm y, \pm z$). This is useful for cross-checking the optical depth structure seen by the observer.

Smoke test. The examples/amr_sphere_generic/ directory ships sphere_amr_inside_test.in (AMR uniform sphere) and sphere_car_inside_test.in (Cartesian 64^3 counterpart), both with par%inside = 8 and an off-center observer at $(o_x, o_y, o_z) = (0.3, 0, 0)$. These produce equivalent HEALPix maps within the documented AMR octree volume-averaging tolerance.

10. Output Files

Outputs are written either as gzip-compressed FITS files (.fits.gz) or as HDF5 files (.h5). HDF5 is the default; set par%file_format = 'fits' in the namelist to switch. Both formats share the same logical layout (one section per FITS HDU $\leftrightarrow$ one HDF5 group), so the conversion is loss-less and runs in both directions via python/lart_io.py (see §11. for the schema and converter, §15. for the Python helper).

10.1 Primary Output Arrays

Array	FITS extension	Description
Jout	JOUT	Emergent spectrum integrated over all escape directions
Jin	JIN	Injected (intrinsic) source spectrum
Jabs	JABS	Spectrum of photons absorbed by dust
Jmu	JMU	Emergent spectrum binned by $\mu = \cos \theta_z$

Jout is the main science output. It is normalized per unit frequency interval and per unit area of the bounding sphere. In dimensionless units (no par%distance_unit):

$$J_{\text{out}}(x) = \frac{\text{escaped photon weight in bin } x}{n_{\text{ph}} \Delta x 2\pi \pi R^2}, \quad (8)$$

where $R = r_{\text{max}}$ (Cartesian) or $R = L_{\text{box}}/2$ (AMR), and $\Delta x = 1$ in units of the Doppler width.

Jin requires par%save_Jin = .true.

Jabs requires par%save_Jabs = .true.

Jmu requires par%save_Jmu = .true.; see §10.2.

The total flux is conserved: $\int J_{\text{out}} dx + \int J_{\text{abs}} dx = \int J_{\text{in}} dx$ (in the absence of other losses).

10.2 Angular Spectrum Jmu

Setting `par%save_Jmu = .true.` adds a 2-D image extension `EXTNAME='Jmu'` of shape (n_x, n_μ) that records the emergent spectrum binned by the z -component of the escape direction, $\mu \equiv \cos \theta_z = \hat{k}_z$. This is useful for

- verifying isotropy of homogeneous models (every μ bin should match `Jout` within Poisson noise),
- characterizing the angular dependence of escape in non-isotropic models (for example, a slab or rotating halo).

Binning.

- `par%nmu` bins of width $\Delta\mu$ tile the μ range uniformly.
- Default `par%nmu = 11` (odd) places $\mu = 0$ at a bin center, so the perpendicular-escape spectrum is preserved as a single channel. Even `par%nmu` places $\mu = 0$ at a bin edge and cleanly separates the $+z$ and $-z$ hemispheres.
- Range:
 - `par%xyz_symmetry = .true.` folds escapes onto the $+z$ hemisphere, so $\mu \in [0, +1]$ and $\Delta\mu = 1/n_\mu$.
 - Otherwise $\mu \in [-1, +1]$ and $\Delta\mu = 2/n_\mu$.

The header keywords `nmu`, `mu_min`, `dmu` record these values, and full WCS keywords (`CTYPE1='XFREQ'`, `CTYPE2='MU'`, `CRVAL*`, `CDELTA*`) are written to the image header.

Normalization. `Jmu` uses the same denominator as `Jout` multiplied by n_μ , so that

$$\frac{1}{n_\mu} \sum_{i=1}^{n_\mu} J(x; \mu_i) = J_{\text{out}}(x) \quad (9)$$

holds exactly. Equivalently, in a homogeneous, isotropic medium each μ bin reproduces `Jout` up to Poisson fluctuations, providing a cheap regression test for new geometries.

Modes. Available with all three grid modes: Cartesian, AMR octree, and the clumpy medium.

10.3 Peel-off cube units and conversion to Jmu units

When peel-off is enabled (`par%nxim`, `par%nyim` > 0 and `par%save_peeloff_3D = .true.`), each observer's `<base>_obs_NNN.fits.gz` file contains Scattered and Direct 3-D image cubes, plus optionally `Direct0`. These are written by `output_sum_rect.f90` as a *specific intensity per pixel*, normalized so that

$$I_{\text{peel}}(x, y, \nu) = \frac{\text{photon weight in pixel } (x, y) \text{ and bin } \nu}{n_{\text{ph}} \delta\Omega_{\text{pix}} \Delta\nu d_{\text{cm}}^2}, \quad (10)$$

where $\delta\Omega_{\text{pix}} = (\text{CD2_2} \cdot \text{CD3_3}) (\pi/180)^2$ is the per-pixel solid angle, $\Delta\nu$ is `Dxfreq` or `Dwave` (set by `par%intensity_unit`), d_{cm}^2 is the square of the distance unit in cm, and n_{ph} is the photon count.

Conversion to Jmu units. J_{out} and J_{mu} are normalized by $4\pi R^2$ (the source surface area) and the outward solid angle 2π , whereas the peel-off cube is normalized by the per-pixel solid angle and d_{cm}^2 . Summing the cube over pixels and rescaling brings the two onto a common axis:

$$J_{\text{peel}}(\nu) = \left(\sum_{(x,y)} I_{\text{peel}}(x,y,\nu) \right) \cdot \frac{4\pi D^2 \delta\Omega_{\text{pix}}}{A_{\text{surface}} \cdot 2\pi}, \quad (11)$$

where D is DISTANCE (in code units) and A_{surface} is the source surface area in code units, with the geometry-dependent forms mirroring `output_sum_rect.f90`:

Geometry	A_{surface} (code units)	Notes
sphere	$4\pi r_{\text{max}}^2$	default; clumpy medium also uses this
box (rectangle)	$8(x_{\text{max}} y_{\text{max}} + y_{\text{max}} z_{\text{max}} + z_{\text{max}} x_{\text{max}})$	full surface of $[-x_{\text{max}}, x_{\text{max}}] \times \dots$ box
slab (par%xy_periodic)	$2(2x_{\text{max}})(2y_{\text{max}})$	two unbounded boundaries
AMR octree	4π	equivalent to $r_{\text{max}} = 1$ (cf. §7.)

The bin width $\Delta\nu$ cancels because both J_{mu} and I_{peel} are already normalized per bin. At convergence and for an axisymmetric source,

$$J_{\text{peel}}(\nu; \alpha, \beta) \rightarrow J_{\text{mu}}(\nu; \mu = \cos \beta). \quad (12)$$

The Python helper `LaRTOutput.plot_peel_jmu_compare()` (in `python/read_lart.py`) implements Eq. (11) with the surface area chosen automatically from `par%geometry`, `par%rmax`, and the spectrum-header `XMAX/YMAX/ZMAX/XY_PER`. See also `examples/sphere/plot_Jpeel_conversion.ipynb` and `examples/clump_sphere/verify_peel_jmu.ipynb`.

Component decomposition (clumpy medium). For a point source inside a clumpy sphere with covering factor f_{cov} , the fraction of photons that escapes without scattering is $\exp(-f_{\text{cov}})$ (Poisson statistics), and the fraction that scatters at least once is $1 - \exp(-f_{\text{cov}})$. J_{mu} tallies *both* populations. The peel-off cube splits them:

- Scattered. The MC estimator over many scatter events averages efficiently, so

$$\int J_{\text{peel}}^{\text{scatt}} d\nu \approx (1 - e^{-f_{\text{cov}}}) \int J_{\text{mu}} d\nu, \quad (13)$$

and the equality holds across all observer angles within Monte-Carlo noise.

- Direct. For a *point* source, every photon shares the same line of sight to a given observer, so the direct contribution is the binary realization of $\exp(-\tau_{\text{LOS}})$: either ~ 1 (clear LOS) or ~ 0 (blocked by an intervening clump) per observer. J_{mu} , by contrast, is azimuth averaged and captures direct escapes uniformly. Adding the binary Direct and the smooth Scattered therefore only matches J_{mu} *in expectation* – it can deviate substantially per observer at finite photon counts.

For shape comparisons in clumpy runs, prefer the scattered component (`component='scatt'` in `plot_peel_jmu_compare()`) so that Eq. (13) provides a clean, deterministic relation.

10.4 Output Control Parameters

Parameter	Default	Description
par%file_format	'hdf5'	Output container: 'hdf5' ($\rightarrow$.h5) or 'fits' ($\rightarrow$.fits.gz); see §11.
par%out_file	''	Output file name (defaults to <input>.h5; extension is corrected to match par%file_format with a warning)
par%out_merge	.false.	Accumulate into existing output instead of overwriting
par%save_backup	.false.	On merge, copy the prior output to <base>_NN.{h5,fits.gz} before overwriting
par%out_bitpix	0	Storage precision: 0=auto (dynamic-range driven), -32=float32, -64=float64
par%save_Jin	.false.	Save injected spectrum
par%save_Jabs	.false.	Save dust-absorbed spectrum
par%save_Jmu	.false.	Save angular spectrum $J(x; \mu)$ (§10.2)
par%nmu	11	Number of μ bins for Jmu
par%save_all	.false.	Save full 3-D $J(\nu, x, y, z)$ array
par%save_direct0	.false.	Save photon escape direction distributions
par%save_radial_profile	.false.	Save 1-D radial J profile
par%save_sightline_tau	.false.	Save optical depth along sightlines
par%intensity_unit	''	Physical unit label written to the FITS header

par%out_merge = .true. is useful for noise reduction: run the same setup multiple times and merge the results. All input parameters must be identical across runs.

11. File Formats: FITS and HDF5

LaRT writes one logical layout that lives in either FITS or HDF5 files. The choice is made by par%file_format (default 'hdf5'); the filename extension is corrected to match if the user-supplied par%out_file disagrees.

11.1 Format selection

- par%file_format = 'hdf5' (default) writes .h5 files via the hdf5io_mod backend. Build with HDF5=1 (the Makefile default).
- par%file_format = 'fits' writes .fits.gz files via the legacy fitsio_mod backend (CFITSIO). Works regardless of the HDF5 build option.

The mismatch warning is issued once on rank 0, e.g. if par%out_file = 'foo.fits.gz' is given alongside par%file_format = 'hdf5':

```
WARNING: par%out_file (foo.fits.gz) does not match
       par%file_format = hdf5
       Overriding par%out_file to: foo.h5
       (set par%file_format explicitly to suppress this warning.)
```

All derived filenames (peel-off _obs, sight-line _tau, _allph, backup) follow par%file_format as well, so the main output and its derived files always share a format.

11.2 HDF5 schema

Each FITS HDU maps to one HDF5 group; the HDU's header keywords become attributes on that group; the HDU's data is either a single /<name>/data dataset (image-style) or one dataset per BinTable column (table-style). Example for a sphere run with peel-off:

```
run.h5
  /Spectrum                                     (group, table-style)
    /Spectrum/Xfreq                           (1-D dataset, float32)
    /Spectrum/velocity                       (1-D dataset, float32)
    /Spectrum/wavelength                     (1-D dataset, float64)
    /Spectrum/Jout                          (1-D dataset, float32)
    /Spectrum/Jin                           (1-D dataset, float32)
    /Spectrum/@nphotons                      (attribute)
    /Spectrum/@taumax                       (attribute)
    /Spectrum/@Nsc_gas                      (attribute)
    ...

run_obs.h5
  /Scattered                                 (group, image-style)
    /Scattered/data                         (3-D dataset, chunked + gzip)
    /Scattered/@CD1_1                      (WCS attribute)
    ...
  /Direct
    /Direct/data
    /Direct/@EXTNAME                       = 'Direct'

run_tau.h5
  /TAU_gas/data, /N_gas/data, [/TAU_dust/data when DGR > 0]
```

Sections appear in their write order on both the root group and inside each subgroup (creation-order tracking is enabled), so h5py iteration returns columns and HDUs in the same order LaRT wrote them. Datasets larger than 4096 elements are chunked and compressed with gzip(4); smaller ones are stored contiguous. The on-disk storage type follows par%out_bitpix (auto / float32 / float64) just as for FITS.

11.3 Conversion utility

The Python helper `python/lart_io.py` reads either format and can convert between them losslessly:

```
python lart_io.py info run.h5
python lart_io.py convert run.fits.gz run.h5
python lart_io.py convert run.h5 run.fits.gz
```

The round-trip preserves all data bit-identically; only string keyword names are uppercased on the FITS leg, matching the FITS card-name convention.

11.4 HDF5 library requirements

LaRT links against the HDF5 Fortran library (`libhdf5_fortran.a` plus the C core `libhdf5.a`, version ≥ 1.14). The Fortran `.mod` files are compiler-specific: an HDF5 build linked against `gfortran` cannot be used with Intel `ifx` or `ifort`, and vice versa. Set `par%HDF5_PREFIX` on the make command line to choose a build that matches the LaRT compiler.

12. Dust and Polarization

Parameter	Default	Description
<code>par%DGR</code>	0.0	Dust-to-gas ratio relative to the Milky Way value
<code>par%albedo</code>	0.32536	Single-scattering albedo ω
<code>par%hgg</code>	0.67617	Heney–Greenstein asymmetry parameter g
<code>par%use_stokes</code>	<code>.false.</code>	Enable full Stokes polarization RT
<code>par%scatt_mat_file</code>	<code>''</code>	Mueller matrix file for polarized dust scattering
<code>par%use_reduced_wgt</code>	<code>.false.</code>	Reduce photon weight by albedo instead of destroying absorbed photons

When `par%DGR = 0` (default), there is no dust and all photon weight is conserved. Non-zero `par%DGR` adds dust extinction and Heney–Greenstein scattering. When `par%use_stokes = .true.`, a Mueller matrix file must be provided; files for the Weingartner–Draine Milky Way model are included in the `data/` directory.

13. Parallelism

Parameter	Default	Description
<code>par%no_photons</code>	10^6	Total photon packets (floating-point, converted to int64)
<code>par%use_master_slave</code>	<code>.true.</code>	Dynamic load balancing via master–worker algorithm
<code>par%num_send_at_once</code>	100	Photons dispatched to each worker per batch
<code>par%nprint</code>	10^5	Progress-report interval [photons]
<code>par%iseed</code>	auto	Random seed; defaults derived from MPI rank and system clock

With `par%use_master_slave = .true.`, one MPI rank acts as dispatcher and the rest as workers. This balances workloads when scattering rates vary across photon packets. With `.false.`, each rank processes an equal share of `par%no_photons`.

14. Complete Parameter Reference

Table 1 lists commonly used parameters. Parameters marked [†] are new or changed in v2.10.

Table 1: LaRT parameter reference.

Parameter	Default	Description
<i>Simulation control</i>		
<code>par%no_photons</code>	10^6	Number of photon packets
<code>par%nprint</code>	10^5	Progress interval [photons]
<code>par%iseed</code>	auto	Random seed
<i>Cartesian grid</i>		
<code>par%nx, par%ny, par%nz</code>	101	Cells per axis
<code>par%xmax, par%ymax, par%zmax</code>	1.0	Half-box size (code units)
<code>par%rmax</code>	-999	Spherical boundary radius
<code>par%rmin</code>	-999	Inner cavity radius
<code>par%geometry</code>	'sphere'	<i>J</i> output geometry (§4.2)
<code>par%xyz_symmetry</code>	.false.	+x, +y, +z octant only
<code>par%xy_symmetry</code>	.false.	+x, +y quadrant only
<code>par%xy_periodic</code>	.false.	Periodic in <i>x, y</i>
<code>par%distance_unit</code>	' '	Physical unit ('kpc', 'pc', 'au')
<code>par%distance2cm</code>	1.0	1 code unit in cm
<i>Grid mode</i>		
<code>par%grid_type[†]</code>	'car'	v2.10+: 'car', 'amr', 'clump'
<code>par%use_amr_grid</code>	.false.	v2.00 AMR flag
<code>par%use_clump_medium</code>	.false.	v2.00 clump flag
<i>Source</i>		
<code>par%source_geometry</code>	'point'	See §5.1
<code>par%xs_point, par%ys_point,</code> <code>par%zs_point</code>	0.0	Point source position
<code>par%source_rmax</code>	-999	Source truncation radius
<code>par%source_rscale</code>	-999	Radial scale length

continued on next page

(continued)

Parameter	Default	Description
par%source_zscale	0.0	Vertical scale height
par%sersic_m	4.0	Sérsic index
par%Reff	1.0	Effective radius
par%comoving_source	.true.	Source co-moving with local gas
par%emiss_file	''	Emissivity FITS cube
par%star_file	''	Point-star list (x, y, z, L)
<i>Line and spectrum</i>		
par%line_id	Ly α	Atomic transition (§5.3)
par%spectral_type	'monochromatic'	Emission frequency distribution
par%temperature	10 ⁴ K	Gas temperature
par%temperature0	= T	Source temperature for voigt0
par%bturb	0.0	Turbulent dispersion [km s ⁻¹]
par%gaussian_sigma_vel	12.84	Gaussian σ [km s ⁻¹] for spectral_type='gaussian'
par%gaussian_FWHM_vel	-999	Gaussian FWHM [km s ⁻¹]; overrides σ if > 0
par%EW_line	0.0	Equivalent width [Å] for 'continuum+gaussian'
par%continuum_normalize	.true.	Normalize output so continuum level = 1
par%xfreq0	0.0	Frequency offset for emitted photons
par%line_prof_file	''	Custom line profile file
par%line_prof_file_type	0	0: (x, I); 1: (λ, I)
<i>Frequency grid</i>		
par%nxfreq	121	Spectral bins
par%xfreq_min, par%xfreq_max	auto	x -grid limits
par%nwavelength	0	Wavelength bins (overrides x -grid)
par%wavelength_min, par%wavelength_max	—	[Å]
par%nvelocity	0	Velocity bins (overrides x -grid)
par%velocity_min, par%velocity_max	—	[km s ⁻¹]

continued on next page

(continued)

Parameter	Default	Description
<i>Optical depth and density</i>		
par%taumax	—	Line-center τ_0 : center $\rightarrow +z$ boundary
par%tauhomo	—	Same, uniform-density equivalent
par%N_HI_max	—	Column density: center $\rightarrow +z$ [cm ⁻²]
par%N_HI_homo	—	Same, uniform-density equivalent
par%N_gas_max	—	Ion column density [cm ⁻²] (metal lines with <code>ion_model='none'</code>)
par%density_rscale	-999	Exponential density scale radius
par%density_alpha	0.0	Power-law density index: $n(r) \propto (r_{\max}/r)^{\alpha_D}$
<i>External data</i>		
par%dens_file	''	Density FITS cube [cm ⁻³]
par%temp_file	''	Temperature FITS cube [K]
par%velo_file	''	Velocity FITS cube [km s ⁻¹]
par%input_field	''	Base name for .dens/.temp/.velo.fits.gz
<i>Velocity</i>		
par%velocity_type	''	Velocity law (§6.4)
par%Vexp	0.0	Expansion/infall speed [km s ⁻¹]
par%velocity_alpha	1.0	Power-law velocity exponent for 'power_law'
par%Vx, par%Vy, par%Vz	0.0	Uniform velocity [km s ⁻¹]
par%Vpeak	0.0	Peak velocity for SSH [km s ⁻¹]
par%rpeak	0.0	Radius at peak for SSH
par%DeltaV	0.0	Velocity increment for SSH [km s ⁻¹]
par%Vrot	0.0	Rotation speed [km s ⁻¹]
par%recoil	.false.	Photon recoil in scattering
<i>AMR</i>		
par%amr_type	'ramses'	'generic' or 'ramses'
par%amr_file	''	AMR data path

continued on next page

(continued)

Parameter	Default	Description
par%amr_snapnum	1	RAMSES snapshot number
<i>Clumpy medium</i>		
par%clump_radius	—	Clump radius
par%clump_tau0	—	Line-center τ_0 per clump
par%clump_f_cov	—	Covering factor C_f
par%clump_f_vol	—	Volume filling factor f_v
par%clump_N_clumps	—	Number of clumps N
par%clump_sigma_v	0.0	Random velocity dispersion [km s ⁻¹]
par%save_clump_info	.false.	Write clump table to FITS
<i>Observers and peel-off</i>		
par%nxim, par%nyim	0	Image pixels; peel-off active when > 0
par%dxim, par%dyim	auto	Pixel size [deg]
par%distance	100 $r_{\max}$	Observer–center distance
par%obsx, par%obsy, par%obsz	(0, 0, 1)	Observer direction
par%alpha, par%beta	—	Azimuth, elevation [deg]
par%gamma	0.0	Detector-plane rotation [deg]
par%save_peeloff	.false.	1-D peel-off spectra
par%save_peeloff_2D	.false.	Peel-off spectral images
par%save_peeloff_3D	.false.	Full 3-D peel-off cube
par%save_direct0	.false.	Peel-off image without $e^{-\tau}$ (transparent-medium reference)
<i>Output</i>		
par%file_format	'hdf5'	Output container: 'hdf5' or 'fits'
par%out_file	''	Output file name (extension follows par%file_format)
par%out_merge	.false.	Accumulate into existing file
par%save_backup	.false.	Keep _NN copy of prior output on merge
par%out_bitpix	0	Storage precision: 0 = auto, -32, -64

continued on next page

(continued)

Parameter	Default	Description
par%save_Jin	.false.	Save injected spectrum
par%save_Jabs	.false.	Save absorbed spectrum
par%save_Jmu	.false.	Save angular spectrum $J(x; \mu)$
par%nmu	11	Number of μ bins for Jmu
par%save_all	.false.	Save full 3-D $J(\nu, x, y, z)$
par%save_radial_profile	.false.	Save 1-D radial J
par%save_sightline_tau	.false.	Save sightline τ
par%intensity_unit	''	Output header unit label
<i>Dust</i>		
par%DGR	0.0	Dust-to-gas ratio
par%albedo	0.32536	Single-scattering albedo
par%hgg	0.67617	Heney–Greenstein asymmetry g
par%use_stokes	.false.	Polarization RT
par%scatt_mat_file	''	Mueller matrix file
par%use_reduced_wgt	.false.	Reduce weight instead of destroy
<i>Parallelism</i>		
par%use_master_slave	.true.	Dynamic load balancing
par%num_send_at_once	100	Photons per worker batch

15. Python Utilities

A small set of post-processing scripts is shipped under `python/` in each version's source tree. They depend on `numpy`, `astropy` (for FITS inputs) and/or `h5py` (for HDF5 inputs), and (for plotting scripts) `matplotlib`. All scripts read both formats transparently via the `lart_io.py` backend (§15.2).

15.1 Loading and analysing output: `read_lart.py`

`read_lart.py` loads a LaRT spectral output (FITS or HDF5) into a rich `LaRTOutput` dataclass that aggregates the angle-integrated spectrum, the optional $J(\mu, x)$ output, and every peel-off observer file alongside it. The dataclass exposes a small library of plotting and analysis methods so common diagnostics can be produced in a few lines. The script accepts the input-file name in any of four forms:

- full input file: `t1tau2.in`
- stem only: `t1tau2` (`.in` is added automatically)
- the HDF5 file: `t1tau2.h5` (also `.hdf5`)

- the FITS file: `t1tau2.fits.gz` (also `.fits`)

When given an input file, the reader honors `par%out_file` from the namelist first, then searches for any LaRT output stem under all recognised extensions (HDF5 preferred) via `lart_io.find_lart_file`. If the output file is still absent but the namelist sets `par%clump_input_file` pointing at an existing clump file, `read_lart` prints a one-line notice and returns a “clumps-only” `LaRTOutput` with the spectrum arrays unset (`out.is_clumps_only == True`) but with the clump population already loaded into `out.clumps` (a `ClumpsOutput`, see below). `out.plot_clump_slice()` therefore works on the pre-run dataclass, so the clump population can be inspected before launching the radiative-transfer run. Peel-off observer files are auto-discovered the same way — both the single-observer (`<stem>_obs.h5`) and the multi-observer (`<stem>_obs_000.h5`, `_obs_001.h5`, ...) layouts are supported, in either format.

Module API.

```

1 from read_lart import read_lart
2 out = read_lart('t1tau2.in')      # FITS or HDF5 -- format-agnostic
3
4 # 1-D arrays, length nxfreq
5 out.xfreq, out.velocity, out.wavelength
6 out.Jout                          # always present
7 out.Jin, out.Jabs, out.Jabs2      # optional -- None if not in file
8
9 # Jmu output -- present only if par%save_Jmu = .true. was used
10 out.mu                          # bin centers, length nm_u
11 out.mu_edges                    # bin edges, length nm_u+1
12 out.Jmu                         # shape (nm_u, nxfreq); None if absent
13 out.nm_u, out.mu_min, out.dmu
14
15 # Per-observer peel-off data (one PeelObservation per observer file)
16 out.peelings                    # list, sorted by (beta, alpha, gamma)
17
18 # Raw section attribute dicts (case-insensitive; legacy uppercase
19 # keys like spec_hdr['EXTNAME'] also work)
20 out.spectrum_header, out.jmu_header

```

Every optional array is `None` when its section is not present, so the reader works on outputs from any LaRT version and either format. The `spectrum_header` and `jmu_header` dicts are filled with both the original-case keys (as they appear in the underlying file) and uppercase aliases, so legacy callers using FITS-style key access (e.g. `spec_hdr['EXTNAME']`) keep working on HDF5 outputs.

Peel-off observer accessor. Each entry in `out.peelings` is a `PeelObservation` dataclass with fields `alpha`, `beta`, `gamma`, `distance`, `nphotons`, `nxim`, `nyim`, `sr_pix`, the scattered + direct cubes `scatt`, `direc`, `direc0` (each shape `(nyim, nxim, nxfreq)`), and the combined cube `= scatt + direc`. Two helper methods are provided:

```

1 po = out.peelings[0]

```

```

1 spec = po.average_spectrum()          # area-mean spectrum, length nxfreq
2 v0, vsig = po.velocity_moment_map(weight='cube')  # (nyim, nxim) each

```

Plotting Jmu. `LaRTOutput.plot_jmu()` draws the angular spectrum directly:

```

1 out = read_lart('t1tau3.in')
2 out.plot_jmu(kind='lines', x='velocity')  # one curve per mu bin
3 out.plot_jmu(kind='image', x='velocity')  # 2-D heatmap (x vs mu)

```

The `x=` keyword selects the abscissa ('velocity', 'xfreq', or 'wavelength'); the `kind=` keyword toggles between line ensemble (colored by μ) and a 2-D image. The line view also overplots `Jout` as a black dashed reference. Pass `ax=` to draw onto an existing axes, `savefig='out.png'` to save to disk. Axis limits accept either the tuple form `xlim=(lo,hi)` / `ylim=(lo,hi)` or the per-bound shorthand `xmin/xmax/ymin/ymax`; the per-bound override takes precedence when both are given:

```

1 out.plot_jmu(kind='lines', xlim=(-500, 500), ymin=0)
2 out.plot_jmu(kind='lines', xmax=300)          # left edge auto-scaled

```

Peel-off analysis plots. Several methods operate on the auto-discovered `out.peelings` list and produce moment maps, per-observer spectra, and radial profiles:

```

1 out.plot_peeling_map(weight='cube', mode='intensity')
2 out.plot_peeling_map(weight='velocity', mode='v_los')
3 out.plot_peeling_map(weight='velocity', mode='sigma_v')
4 out.plot_peeling_spectrum()          # one curve per observer
5 out.plot_peeling_radial_profile()     # azimuthally averaged
6 out.plot_velocity_moment_map(po, mode='v_los')  # single observer

```

The `plot_peel_jmu_compare()` method overlays each observer's peel-off spectrum on the matching $J(\mu, x)$ slice as an isotropy diagnostic, and `plot_clump_slice()` draws the 2D cross-section of the clumpy medium at a chosen plane (it opens the matching `<stem>.clumps.h5` / `.fits.gz` file automatically). All plotting methods accept `ax=`, `savefig=`, and `xlim/ylim` (or `xmin/xmax/ymin/ymax`) the same way as `plot_jmu`.

Command-line usage.

```

1 python read_lart.py t1tau2          # auto-resolves to t1tau2.in
2 python read_lart.py t1tau2.in
3 python read_lart.py t1tau2.h5       # HDF5 output
4 python read_lart.py t1tau2.fits.gz  # FITS output

```

prints a one-screen summary of which arrays are present and, when `Jmu` is in the file, the per-bin $\max_x J(x; \mu) / \max_x J_{\text{out}}(x)$ ratio (a quick isotropy sanity check for homogeneous models).

Standalone clump access: `read_clumps` / `read_clump`. The clump population can also be loaded independently of any radiative-transfer output via the top-level helper:

```

1 from read_lart import read_clumps      # alias: read_clump
2 co = read_clumps('foo_clumps.h5')     # direct clump file
3 co = read_clumps('run.in')            # via par%clump_input_file
4 co = read_clumps('run')                # stem -> tries .in, then _clumps.*
5 print(co.summary())
6 co.plot_clump_slice(axis='z')

```

The returned ClumpsOutput dataclass exposes the per-clump arrays (x, y, z, vx, vy, vz, radius, rhokap, density, temp, plus the convenience attributes pos, vel, n_clumps, sphere_r, f_vol, f_cov, distance_unit, distance2cm), the case-insensitive attribute dict attrs (with co.attr(name) helper), and the plot_clump_slice() method. The opacity column in the file is auto-detected with priority RHOKAP → DENSITY → DENS: RHOKAP populates co.rhokap (opacity per code unit), while DENSITY / DENS populate co.density (n_{HI} [cm⁻³]); co.opacity_col records which column was found. co.distance_unit / co.distance2cm expose the DISTUNIT / DIST_CM attributes from the file (LaRT's par%distance_unit / par%distance2cm at write time). When the F_VOL / F_COV attributes are missing from the file, the f_vol / f_cov properties fall back to compute_f_vol() / compute_f_cov(), which evaluate the same formulas as LaRT's write_clumps_fits() from-file branch in clump_mod.f90: $f_{vol} = \sum r_i^3 / (R^3 - r_{min}^3)$ and $f_{cov} = \frac{3}{4} \sum r_i^2 / (R^2 + R r_{min} + r_{min}^2)$, with R taken from SPHERE_R (or RMAX) and r_{min} from RMIN (default 0). The two compute_* methods can also be called explicitly to recompute regardless of the stored attribute. Inside LaRTOutput the same dataclass lives at out.clumps (lazy- loaded on the first plot_clump_slice call, or populated eagerly in the clumps-only pre-run state described above).

Other scripts. The following plotting helpers are also shipped:

- plot_spectrum.py, plot_map.py, plot_intensity_profile.py, plot_Pa.py: standard quick-look plots.
- check_flux.py, flux_check.py: flux-conservation diagnostics.
- AMR_grid/: helpers for generating generic AMR data files (see AMR_grid/make_amr_grid.py, make_amr_sphere_radial.py).

15.2 Format-agnostic reader and FITS ↔ HDF5 converter: lart_io.py

python/lart_io.py provides a single load_lart(path) entry point that reads either a FITS or an HDF5 LaRT output and returns a uniform LartFile structure (one Section per HDU / group, with .data or .columns plus .attrs). It also converts between the two formats by extension. Requires numpy plus astropy (for FITS) and/or h5py (for HDF5); each backend is only imported when needed.

Module API.

```

1 from lart_io import load_lart, convert
2
3 f = load_lart('run.h5')              # works on .fits.gz too
4 spec = f['Spectrum']                 # a Section

```



```

5 jout = spec.columns['Jout']          # 1-D numpy array
6 nph  = spec.attrs['nphotons']
7
8 scat = load_lart('run_obs.h5')['Scattered']
9 cube = scat.data                    # 3-D array (nxim, nyim, nxfreq)
10
11 convert('run.fits.gz', 'run.h5') # both directions by extension
12 convert('run.h5',          'run.fits.gz')

```

The round-trip preserves all numerical data bit-identically; string keyword names are normalized to uppercase on the FITS leg (FITS card-name convention), which is the only visible difference between the two forms.

Command-line usage.

```

1 python lart_io.py info  run.h5
2 python lart_io.py info  run.fits.gz
3 python lart_io.py convert run.fits.gz run.h5
4 python lart_io.py convert run.h5      run.fits.gz

```

16. Example Input Files

16.1 Uniform Sphere, Ly α , Static

```

1 &parameters
2   par%no_photons      = 1e6
3   par%temperature     = 1.0e4
4   par%taumax          = 1.0e4
5   par%rmax            = 1.0
6   par%nx = 101, par%ny = 101, par%nz = 101
7   par%source_geometry = 'point'
8   par%spectral_type   = 'voigt'
9   par%nxfreq          = 121
10 /

```

16.2 Sphere with Hubble Outflow and Peel-off

```

1 &parameters
2   par%no_photons      = 1e6
3   par%temperature     = 1.0e4
4   par%taumax          = 1.0e4
5   par%rmax            = 1.0
6   par%nx = 101, par%ny = 101, par%nz = 101

```

```

7 par%source_geometry = 'point'
8 par%spectral_type   = 'voigt'
9 par%velocity_type   = 'hubble'
10 par%Vexp            = 200.0          ! km/s outflow
11 par%nxfreq          = 121
12 par%nxim = 129, par%nyim = 129
13 par%alpha = 0.0 45.0 90.0
14 par%beta  = 0.0 0.0 0.0
15 par%save_peeloff_2D = .true.
16 /

```

16.3 Realistic Galaxy (FITS Density Field, Stokes Polarization)

```

1 &parameters
2 par%no_photons    = 1e8
3 par%dens_file     = 'galaxy.dens.fits.gz'
4 par%temp_file     = 'galaxy.temp.fits.gz'
5 par%velo_file     = 'galaxy.velo.fits.gz'
6 par%taumax        = 1.0e4           ! rescale to this optical depth
7 par%distance_unit = 'kpc'
8 par%source_geometry = 'exponential'
9 par%source_rscale = 3.0
10 par%source_zscale = 0.5
11 par%spectral_type = 'voigt'
12 par%DGR           = 1.0
13 par%use_stokes    = .true.
14 par%nxfreq        = 121
15 par%nxim = 200, par%nyim = 200
16 par%save_peeloff_2D = .true.
17 /

```

16.4 AMR Sphere (Generic Text Format)

```

1 &parameters
2 par%use_amr_grid  = .true.          ! or par%grid_type = 'amr' in v2.10
3 par%amr_type      = 'generic'
4 par%amr_file      = 'sphere_amr.dat'
5 par%no_photons    = 1e6
6 par%taumax        = 1.0e4
7 par%source_geometry = 'point'

```

```

8  par%spectral_type = 'voigt'
9  par%nxfreq       = 121
10 par%nxim = 100, par%nyim = 100
11 par%save_peeloff_2D = .true.
12 par%save_Jin       = .true.
13 /

```

16.5 Clumpy Sphere with Random Velocities

```

1 &parameters
2  par%use_clump_medium = .true.      ! or par%grid_type = 'clump' in v2.10
3  par%rmax             = 1.0
4  par%clump_radius     = 0.01
5  par%clump_f_cov      = 5.0
6  par%clump_tau0       = 1.0e3
7  par%temperature      = 1.0e4
8  par%clump_sigma_v    = 50.0       ! km/s random dispersion
9  par%no_photons       = 1e6
10 par%spectral_type     = 'voigt'
11 par%nxfreq           = 301
12 par%velocity_min     = -200.0
13 par%velocity_max     = 200.0
14 par%save_clump_info  = .true.
15 /

```

16.6 Illustris/TNG Galaxy (Diffuse Emissivity)

After converting a TNG cutout with `convert_illustris_to_generic.py --compute-physics`:

```

1 &parameters
2  par%use_amr_grid     = .true.
3  par%amr_type         = 'generic'
4  par%amr_file         = 'galaxy.h5'
5  par%distance_unit    = 'kpc'
6  par%no_photons       = 1e6
7  par%spectral_type    = 'voigt'
8  par%source_geometry  = 'diffuse_emissivity'
9  par%core_skip        = .true.
10 par%nxfreq           = 201
11 par%velocity_min     = -1000.0
12 par%velocity_max     = 1000.0

```

```

13 par%save_peeloff      = .true.
14 par%nxim = 128, par%nyim = 128
15 par%distance          = 1e4
16 par%alpha = -90.0
17 par%beta  = 0.0
18 /

```

The emissivity column in galaxy.h5 is automatically used as the photon source distribution. No par%emiss_file or par%emissivity_model is required.

Metal-line Examples (C IV, O VI)

The converter supports CIE collisional excitation emissivity for metal lines via --emissivity-line:

```

1 # C IV 1548+1550
2 python convert_illustris_to_generic.py cutout.hdf5 \
3     -o galaxy_civ.h5 --compute-physics --emissivity-line civ
4
5 # O VI 1032+1038
6 python convert_illustris_to_generic.py cutout.hdf5 \
7     -o galaxy_ovi.h5 --compute-physics --emissivity-line ovi

```

Then run LaRT with the appropriate par%line_id:

```

1 &parameters
2 par%use_amr_grid      = .true.
3 par%amr_type          = 'generic'
4 par%amr_file          = 'galaxy_ovi.h5'
5 par%distance_unit     = 'kpc'
6 par%line_id           = 'OVI_1032'
7 par%no_photons        = 1e6
8 par%spectral_type     = 'voigt'
9 par%source_geometry   = 'diffuse_emissivity'
10 par%core_skip         = .true.
11 par%nxfreq            = 201
12 par%velocity_min      = -1500.0
13 par%velocity_max      = 2500.0
14 /

```

Note: the O VI doublet separation is $\sim 1652 \text{ km s}^{-1}$ (vs. $\sim 500 \text{ km s}^{-1}$ for C IV), so a wider velocity range is needed.